

# Neural Network Assignment

Nilanjan<sup>a,\*</sup>

<sup>a</sup>Department of Molecular Biosciences, University of Kansas

\*Corresponding author, nilanjan.roy@ku.edu

November 28, 2024

## 1 Data

For the complex dataset, I used a MRI image dataset for classification of 3 different tumors. The dataset has 4 folders consisting of images of glioma, meningioma, pituitary tumor along with normal brain MRI images. To reduce computational load, I chose to use 200 images of each of these categories. The link to download the dataset is given in a .txt file in the data folder.

## 2 Day 1: Building the first NN model

Model 1: Here, I used a simple model, with 2 dense and 1 dropout layer. The input shape was comparatively big (196608) which is why I decided to go with 3 layers as it was taking a huge time to train. Model 1 Code:

```
# Model 1: Simple Feedforward Neural Network
model1 <- keras_model_sequential() %>%

# First dense layer with 64 units and ReLU activation
layer_dense(units = 64, activation = 'relu', input_shape = c(196608)) %>%

# Dropout layer to prevent overfitting (20% dropout rate)
layer_dropout(rate = 0.2) %>%

# Output layer with 4 units (one for each class) and softmax activation
layer_dense(units = 4, activation = 'softmax')

# Display the model architecture
summary(model1)
# Compile the model
model1 %>% compile(
  loss = 'sparse_categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = c('accuracy')
)

history1 <- model1 %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 128,
```

```
validation_split = 0.2
)
```

Model 1 evaluation table:

Metric	Value (Train)	Value (Validation)
Accuracy	0.2988	0.3047
Loss	1.327	1.393

Table 1: Day1-Model1: Final Epoch Metrics: Training and Validation

Model 1 accuracy visualization:




Figure 1: Day1-Model1: Accuracy visualization

The accuracy is around 0.30 in this model. That is expected as the input shape size is huge and other were only 3 layers. Or this can be due to low amount of images (200) for each category.

### 3 Day 2: Building additional NN models on the complex data

Model 2 and 3: In model 2 and 3 I increased the amount of layers with different parameters. Model 2 and 3 Code:

```
# Model 2: Adding Another Dense Layer and Dropout
model2 <- keras_model_sequential() %>%

# First dense layer with 64 units and ReLU activation
layer_dense(units = 64, activation = 'relu', input_shape = c(196608)) %>%

# Dropout layer with 60% dropout rate
layer_dropout(rate = 0.6) %>%
```

```

# Second dense layer with 32 units and ReLU activation
layer_dense(units = 32, activation = 'relu') %>%

# Dropout layer with 40% dropout rate
layer_dropout(rate = 0.4) %>%

# Output layer with 4 units and softmax activation
layer_dense(units = 4, activation = 'softmax')

# Display the model architecture
summary(model2)

# Model 3: Adjusting Sizes and Dropout Rates
model3 <- keras_model_sequential() %>%

# First dense layer with 128 units and ReLU activation
layer_dense(units = 128, activation = 'relu', input_shape = c(196608)) %>%

# Dropout layer with 50% dropout rate
layer_dropout(rate = 0.5) %>%

# Second dense layer with 64 units and ReLU activation
layer_dense(units = 64, activation = 'relu') %>%

# Dropout layer with 30% dropout rate
layer_dropout(rate = 0.3) %>%

# Output layer with 4 units and softmax activation
layer_dense(units = 4, activation = 'softmax')

# Display the model architecture
summary(model3)

# Compile both models
model2 %>% compile(
  loss = 'sparse_categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = c('accuracy')
)

model3 %>% compile(
  loss = 'sparse_categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = c('accuracy')
)

history2 <- model2 %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 128,
  validation_split = 0.2
)

```

```
history3 <- model3 %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 128,
  validation_split = 0.2
)
```

Model 2 evaluation table:

Metric	Value (Train)	Value (Validation)
Accuracy	0.2637	0.2422
Loss	1.38	1.376

Table 2: Day2-Model2: Final Epoch Metrics: Training and Validation

Model 3 evaluation table:

Metric	Value (Train)	Value (Validation)
Accuracy	0.2578	0.2734
Loss	1.361	1.395

Table 3: Day2-Model3: Final Epoch Metrics: Training and Validation

Model 2 accuracy visualization:




Figure 2: Day2-Model2: Accuracy visualization

Model 3 accuracy visualization:




Figure 3: Day2-Model3: Accuracy visualization

So, even after increasing the dense layers the accuracy was horrible which is around 0.30 for these three models that were tried in this dataset. Here, I did not build the function to look at the misclassification as all of the models are equally bad. But it gets better in Day3 when we applied CNN. I used the misclassification function to inspect some of the misclassifications there.

## 4 Day 3: Applying CNN and different optimizers

3 different CNN Model on the MRI data:

```
# Model 1: Convolutional Neural Network with Dropout Layers
model1 <- keras_model_sequential() %>%
  layer_conv_2d(filters = 16, kernel_size = c(3, 3), padding = "same",
    input_shape = c(img_height, img_width, num_channels), activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(4, 4)) %>%
  layer_dropout(rate = 0.4) %>%
  layer_conv_2d(filters = 32, kernel_size = c(3, 3), padding = "same",
    activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_dropout(rate = 0.8) %>%
  layer_conv_2d(filters = 16, kernel_size = c(3, 3), padding = "same",
    activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_dropout(rate = 0.2) %>%
  layer_flatten() %>%
  layer_dense(units = 64, activation = "relu") %>%
  layer_dense(units = 4, activation = "softmax")

# Display the model architecture
summary(model1)
```

```

# Compile the model
model1 %>% compile(
  loss = 'sparse_categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = c('accuracy')
)

# Train the Model

history1 <- model1 %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 32,
  validation_split = 0.2
)

# View training history
history1
plot(history1)

# Build Alternative Models

# Model 2: Adjusted Dropout Rates
model2 <- keras_model_sequential() %>%
  layer_conv_2d(filters = 16, kernel_size = c(3, 3), padding = "same",
    input_shape = c(img_height, img_width, num_channels), activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(4, 4)) %>%
  layer_dropout(rate = 0.8) %>%
  layer_conv_2d(filters = 32, kernel_size = c(3, 3), padding = "same",
    activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_dropout(rate = 0.8) %>%
  layer_conv_2d(filters = 16, kernel_size = c(3, 3), padding = "same",
    activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_dropout(rate = 0.2) %>%
  layer_flatten() %>%
  layer_dense(units = 64, activation = "relu") %>%
  layer_dense(units = 4, activation = "softmax")

# Display the model architecture
summary(model2)

# Compile the model
model2 %>% compile(
  loss = 'sparse_categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = c('accuracy')
)

# Train Model 2

```

```

history2 <- model2 %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 32,
  validation_split = 0.2
)

# Model 3: Simplified Model with Fewer Layers
model3 <- keras_model_sequential() %>%
  layer_conv_2d(filters = 16, kernel_size = c(3, 3), padding = "same",
    input_shape = c(img_height, img_width, num_channels), activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(4, 4)) %>%
  layer_dropout(rate = 0.8) %>%
  layer_conv_2d(filters = 16, kernel_size = c(3, 3), padding = "same",
    activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_dropout(rate = 0.2) %>%
  layer_flatten() %>%
  layer_dense(units = 64, activation = "relu") %>%
  layer_dense(units = 4, activation = "softmax")

# Display the model architecture
summary(model3)

# Compile the model
model3 %>% compile(
  loss = 'sparse_categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = c('accuracy')
)

# Train Model 3
history3 <- model3 %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 32,
  validation_split = 0.2
)

```

CNN Model evaluation table:

Model	Accuracy (Train)	Loss (Train)	Accuracy (Validation)	Loss (Validation)
CNN Model 1	0.9082	0.2493	0.7891	0.57
CNN Model 2	0.9102	0.2693	0.7891	0.6919
CNN Model 3	0.9883	0.04876	0.7266	1.445

Table 4: Day 3: Final Epoch Metrics for CNN Models

CNN model accuracy visualization:

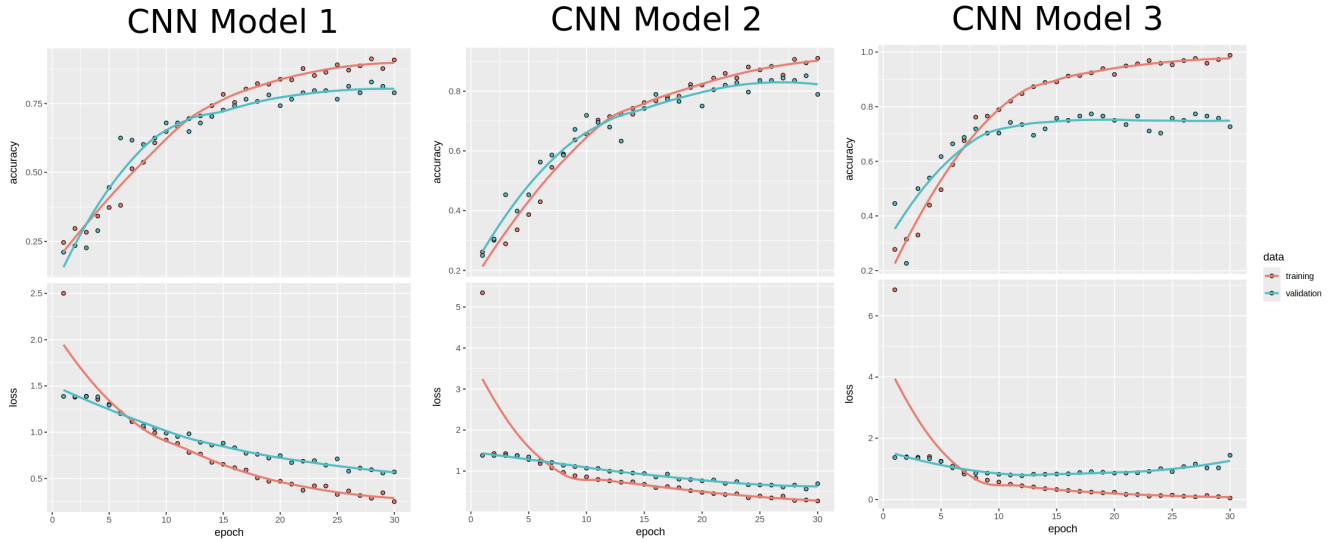


Figure 4: Day3-CNN Models and their Accuracy

After using the CNN models on the same 200 MRI images from each category the accuracy increased tremendously. Model 1 and 2 had similar success. Model 3 accuracy rate was super high but did not do that much well in the validation set. Model 1 and 2 also had some issues with their validation set. Altogether I decided that model 2 might be the best model among these three.

Looking at some misclassified images by CNN model 2:

Model says: meningioma but reality is: normal

Code	Class	Probability
0	Glioma	0.007
1	Meningioma	0.531
2	Normal	0.434
3	Pituitary	0.027

Table 5: Class Probabilities for Different Codes by model 2

Lets see the misclassified image:




Figure 5: Day3-CNN Models miss-classification

Looks like a normal brain image. This can be a model mistake.

Now, I tried 3 different optimizers (optimizer:lamb, optimizer:lion, optimizer:adam) on CNN model 2. Accuracy evaluation of the three different optimizer:



Optimizer	Accuracy (Train)	Loss (Train)	Accuracy (Validation)	Loss (Validation)
Lamb	0.7832	0.5524	0.7188	0.9091
Lion	0.2617	1.381	0.2344	1.409
Adam	0.5703	0.9315	0.5859	1.122

Table 6: Day 3: Epoch Metrics for Different Optimizers on CNN model 2

None of these optimizer made the CNN model 2 better (made things worst). The original optimizer:rmsprop actually did best.

## 5 Day 4: Transfer learning

Here again, I used 200 images of each tumor category to test whether transfer learning improves my previous models.

Scratch CNN model code:

```
# Define the Convolutional Neural Network model
cnn_model <- keras_model_sequential() %>%
  # First convolutional layer
  layer_conv_2d(filters = 16, kernel_size = c(3, 3), padding = "same",
    input_shape = c(img_height, img_width, num_channels), activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(4, 4)) %>%
  layer_dropout(rate = 0.4) %>%
  # Second convolutional layer
  layer_conv_2d(filters = 32, kernel_size = c(3, 3), padding = "same", activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_dropout(rate = 0.3) %>%
  # Third convolutional layer
  layer_conv_2d(filters = 32, kernel_size = c(3, 3), padding = "same", activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_dropout(rate = 0.2) %>%
  # Flatten the output and add dense layers
  layer_flatten() %>%
  layer_dense(units = 64, activation = "relu") %>%
  # Output layer
  layer_dense(units = 4, activation = "softmax")

# Display the model architecture
summary(cnn_model)

# Compile the CNN model
cnn_model %>% compile(
  loss = 'sparse_categorical_crossentropy',
  optimizer = optimizer_adam(),
  metrics = c('accuracy')
)

# Train the CNN Model -----

cnn_history <- cnn_model %>% fit(
  x_train, y_train,
```

```

batch_size = 32,
epochs = 10,
validation_split = 0.2
)

```

Transfer learning model:

```

# Load the pre-trained Xception model without the top layers
base_model <- application_xception(
  weights = "imagenet",      # Load weights pre-trained on ImageNet
  include_top = FALSE,       # Do not include the final output layer
  input_shape = c(img_height, img_width, num_channels) # Adjust input shape
)

# Freeze the base model layers to prevent them from being updated during training
base_model$trainable <- FALSE

# Display the base model architecture
summary(base_model)

# Build the new model on top of the base model
transfer_model <- keras_model_sequential() %>%
  base_model %>%
  # Add global average pooling layer
  layer_global_average_pooling_2d() %>%
  # Add dropout for regularization
  layer_dropout(rate = 0.2) %>%
  # Add output layer
  layer_dense(units = 4, activation = "softmax")

# Display the new model architecture
summary(transfer_model)

# Compile the Transfer Learning Model -----

transfer_model %>% compile(
  loss = 'sparse_categorical_crossentropy',
  optimizer = optimizer_adam(),
  metrics = c('accuracy')
)

# Train the Transfer Learning Model -----

transfer_history <- transfer_model %>% fit(
  x_train, y_train,
  batch_size = 32,
  epochs = 6,
  validation_split = 0.2
)

```

Accuracy evaluation of the scratch CNN model and transfer learning model:

Model	Accuracy (Train)	Loss (Train)	Accuracy (Validation)	Loss (Validation)
Scratch CNN	0.8496	0.4384	0.7812	0.6696
Transfer Learning	0.9375	0.23	0.8359	0.4216

Table 7: Day 4: Final Epoch Metrics: Scratch CNN vs. Transfer Learning Model




Figure 6: Day 4: Scratch Models and transfer learning evaluation

So, the transfer learning improved the scratch CNN. Although the plot of transfer learning looks weird. But if you look at the y-axis scale is between 0.82 to 0.93. It may be over fitted but may be still better than scratch CNN.